

Assignment

Roll No.: CS23B1047

Name: Dhage pratik Bhishmacharya

course: Database systems.

Q.1 compare and contrast B-Trees and B+ trees. Explain their structure and use cases.

Solution:

B Tree structure - A B-Tree is self-balancing, multi-level index designed for efficient disk-based storage. It generalizes binary search trees to allow multiple keys per node and multiple children, minimizing I/O operations.

order m : Nodes have up to $m-1$ keys & m children.

keys: sorted, up to $m-1$, minimum is $\lceil \frac{m}{2} \rceil - 1$ (except root node has minimum 1)

pointers: $K+1$ children for K keys

B+ Tree structure - A B+ tree is a variant of the B tree optimized for range queries and sequential access, commonly used in database indexing.

order m : Internal nodes have up to $m-1$ keys, m children; leaf nodes store up to m keys.

Internal nodes: keys only, pointers to children, No data.

Leaf Nodes: All keys and data, sorted, linked sequentially (linked-list)

Balanced, with minimum $\lceil \frac{m}{2} \rceil$ keys in leaves, $\lceil \frac{m}{2} \rceil$ children in internal nodes. (except root.)

Comparison:

| feature | B-Tree | B+ Tree |
|------------------|-------------------------------------|--------------------------------------|
| Data storage | Internal and leaf nodes | Leaf nodes only |
| Leaf Linking | None | Doubly-linked list |
| search path | May end at internal nodes. | Always end at leaf node. |
| Range queries | Less efficient, requires traversal. | Efficient, sequential leaf access. |
| fanout | Lower (data reduces key capacity) | Higher (keys only in internal nodes) |
| space efficiency | Less efficient | More efficient. |

Use cases:

1) B-Tree -

- (a) File systems: Used in file systems like NTFS and HFS+ for hierarchical, directory structures.
- (b) Database with frequent updates.
- (c) Embedded systems - preferred in memory-constrained environments.

2) B+ Tree -

- (a) Database Indexing: used in Relational DBMS for indexing tables.
- (b) Data Warehouses: Ideal for analytical queries requiring scans over large datasets.
- (c) search engines: supports efficient keyword-based range searches.

Conclusion: B+ Trees are generally preferred in database systems due to their optimization for range queries and sequential access.

Q.2. Explain the insertion operation in a B-tree with an example. How does it maintain balance?

Solution: Insertion in a B-tree maintains balance by ensuring nodes adhere to key limits and the tree remains height balanced. Steps:

(1) Locate leaf: Traverse from root to the appropriate leaf node using key comparisons.

(2) Insert key:

If node has $< m-1$ keys; insert in sorted order

If overflow, $m-1$ keys full, split:

i. Divide into two nodes with $\lfloor \frac{m}{2} \rfloor$ and $\lceil \frac{m}{2} \rceil - 1$ keys.

ii. promote median key to parent

iii. Adjust child pointers

(3) propagate splits: If parent is full, split recursively up to root. If root splits, create new root, increasing height.

(4) Balance: Nodes (except root) maintain $\lfloor \frac{m}{2} \rfloor - 1$ to $m-1$ keys.

Example: create a B-Tree of order 3 by inserting values from 1 to 10.

$$\text{max keys} = m-1 = 3-1 = 2$$

$$\text{min keys} = \lfloor \frac{m}{2} \rfloor - 1 = 2-1 = 1$$


Insert 1: $\boxed{} \boxed{} \boxed{} \rightarrow \boxed{1} \boxed{} \boxed{}$ Insert 2 $\rightarrow \boxed{1} \boxed{2} \boxed{}$

Insert 3 \downarrow


$\boxed{1} \boxed{2} \boxed{3}$

overflow.

Insert 4 \leftarrow




Insert 5 \swarrow




overflow.


Insert 6 \rightarrow



Insert 7



Insert 10



Balance Maintenance:

Node splitting: When a node becomes full, it is split into two nodes, distributing keys evenly & promoting the median key.

Minimum key constraint: After splitting, each node (except root) has at least $\lceil \frac{m}{2} \rceil - 1$ keys.

Height adjustment: Splitting may increase the tree height only when root splits, ensuring logarithmic height.

Recursive propagation: splits propagate upwards, adjusting the tree structure to maintain balance.

Q.3. Describe the deletion process in a B+ tree. What challenges are faced and how are they resolved?

Solution:

Deletion process:

- (1) Locate key: Traverse to leaf node ~~and~~ ~~contains~~ containing the key.
- (2) Delete key:
 - (a) Remove key from leaf
 - (b) If leaf has $\geq \lceil \frac{m}{2} \rceil$ keys, done.
 - (c) If underflow ($< \lceil \frac{m}{2} \rceil$ keys), balance.

(3) Handle underflow:

(a) Borrow: If sibling has $> \lfloor \frac{m}{2} \rfloor$ keys, borrow a key via parent, update parent's separator.

(b) Merge: If sibling has $\leq \lfloor \frac{m}{2} \rfloor$ keys, merge with sibling and parent's separator key, remove separator from parent.

propagate underflow to parent if needed.

(4) Update internal nodes: If key was a separator, replace with predecessor / successor from leaf.

(5) Root adjustment: If root has one child post-deletion, make child the new root, reducing height.

Challenges and Resolutions:

(1) underflow in leaves -

challenge: Deletion leaves node with $< \lfloor \frac{m}{2} \rfloor$ keys.

Resolution: Borrow from sibling or merge, preferring borrowing to maintain height.

(2) parent underflow -

challenge: Merging leaves causes parent to underflow.

Resolution: Borrow or merge at parent level, propagating recursively.

(3) separator key updates -

challenge: Deleted key in internal nodes must be replaced.

Resolution: Use predecessor / successor from leaf to maintain search property.

(4) Height Reduction -

challenge: Root may have one child after merges.

Resolution: Set child as new root, reducing height.

Q4. Discuss the advantages and disadvantages of using B+ Trees over B-trees in database indexing.

Advantages of B+ Trees:

- (1) Efficient Range Queries - linked leaf nodes allow sequential access, making range queries faster than B-Trees.
- (2) Higher fanout - Internal nodes store only keys, not data, allowing more keys per node. This reduces tree height, improving search.
- (3) Simplified search path - All data is at leaf nodes, ensuring consistent search paths to leaves.
- (4) Better space optimization - storing data only in leaves allows internal nodes to hold more keys, reducing storage overhead compared to B-Trees.
- (5) Support for sequential access - critical for database operations like table scans or joins.

Disadvantages of B+ Trees:

- (1) Complex maintenance - maintaining the linked list of leaf nodes adds complexity during insertions & deletions.
- (2) Slower point queries - point queries may be slightly slower than in B-trees since data is only in leaves.
- (3) Higher overhead for updates - Insertions and deletions require updating the linked list & potentially adjusting separator keys in internal nodes, increasing overhead.
- (4) Memory usage for linking - the doubly-linked list in leaf nodes requires additional pointers, increasing memory usage, which may be a concern in memory-constrained systems.

Q5. Construct a B+ Tree of order 4 by inserting the following sequence of keys: 10, 20, 5, 6, 12, 30, 7, 17.
Show each step clearly.

solution: order (m)=4

min. no. of keys = 1 (root), 2 (leaf & internal)

max no. of keys = $4-1=3$ (internal), 2 (leaf)

create and insert 10.


[10]

insert 20 and insert 5.


[5 | 10 | 20]

insert 6. [5 | 6 | 10 | 20] ... overflow (full).


∴ split and promote 10 to parent




insert 12 :



insert 30 :



∴ split and promote 20 to parent.



insert 7 and insert 17

