

SQL

SQL – Structured Query Language

SEQUEL – Structured English QUERy Language

Parts of SQL:

DDL

DML

Views

Security & Authorization

Transaction Control

CREATE SCHEMA

CREATE SCHEMA COMPANY AUTHORIZATION 'Jsmith';

Authorization Identifier – Who owns the Schema

In general, not all users are authorized to create schemas and schema elements. The privilege to create schemas, tables, and other constructs must be explicitly granted to the relevant user accounts by the system administrator or DBA.

Tables are generally created within the SQL Schema

Do you
Find
Something
Weird
In this ?

CREATE TABLE EMPLOYEE

| | | |
|-----------|----------------|-----------|
| (Fname | VARCHAR(15) | NOT NULL, |
| Minit | CHAR, | |
| Lname | VARCHAR(15) | NOT NULL, |
| Ssn | CHAR(9) | NOT NULL, |
| Bdate | DATE, | |
| Address | VARCHAR(30), | |
| Sex | CHAR, | |
| Salary | DECIMAL(10,2), | |
| Super_ssn | CHAR(9), | |
| Dno | INT | NOT NULL, |

PRIMARY KEY (Ssn),

FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn),

FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber) ;

Foreign Keys can be specified later, through ALTER Command

CREATE TABLE EMPLOYEE

**Attributes are considered
To be ordered
In this order**

| | | |
|-----------|----------------|------------------|
| (Fname | VARCHAR(15) | NOT NULL, |
| Minit | CHAR, | |
| Lname | VARCHAR(15) | NOT NULL, |
| Ssn | CHAR(9) | NOT NULL, |
| Bdate | DATE, | |
| Address | VARCHAR(30), | |
| Sex | CHAR, | |
| Salary | DECIMAL(10,2), | |
| Super_ssn | CHAR(9), | |
| Dno | INT | NOT NULL, |

PRIMARY KEY (Ssn),
FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn),
FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber) ;

4.2.2 Specifying Key and Referential Integrity Constraints

Dnumber INT **PRIMARY KEY**;

PRIMARY KEY = NOT NULL + UNIQUE

Is it possible to create a Table in SQL without a Primary Key ?

4.1.3 Attribute Data Types and Domains in SQL

- Numeric data types include integer numbers of various sizes (INTEGER or INT, and SMALLINT) and floating-point (real) numbers of various precision (FLOAT or REAL, and DOUBLE PRECISION). Formatted numbers can be declared by using DECIMAL(*i,j*)—or DEC(*i,j*) or NUMERIC(*i,j*)—where *i*, the *precision*, is the total number of decimal digits and *j*, the *scale*, is the number of digits after the decimal point. The default for scale is zero, and the default for precision is implementation-defined.

4.1.3 Attribute Data Types and Domains in SQL

Character String

Bit String

Boolean

Date

Timestamp

4.1.3 Attribute Data Types and Domains in SQL

```
CREATE DOMAIN SSN_TYPE AS CHAR(9);
```

4.2.1 Specifying Attribute Constraints and Attribute Defaults

NOT NULL

DEFAULT

Dnumber INT **NOT NULL CHECK** (Dnumber > 0 **AND** Dnumber < 21);

**CREATE DOMAIN D_NUM AS INTEGER
CHECK (D_NUM > 0 **AND** D_NUM < 21);**

4.2.2 Specifying Key and Referential Integrity Constraints

Dnumber INT **PRIMARY KEY**;

CREATE TABLE DEPENDENT

| | | |
|----------------|-------------|------------------|
| (Essn | CHAR(9) | NOT NULL, |
| Dependent_name | VARCHAR(15) | NOT NULL, |
| Sex | CHAR, | |
| Bdate | DATE, | |
| Relationship | VARCHAR(8), | |

PRIMARY KEY (Essn, Dependent_name),
FOREIGN KEY (Essn) **REFERENCES** EMPLOYEE(Ssn));

 **Composite Primary Key**

4.2.2 Specifying Key and Referential Integrity Constraints

Dname VARCHAR(15) **UNIQUE**;

Can potentially
Act as a
Secondary / Alternate Key

Foreign Key

Syntax for F.Key you know already.

What about actions when referential integrity is violated ?

The default action is to “REJECT” the update operation.

Referential Triggered Action

```
CREATE TABLE DEPARTMENT
(
    ...,
    Mgr_ssn      CHAR(9)          NOT NULL          DEFAULT '888665555',
    ...
    CONSTRAINT DEPTPK
        PRIMARY KEY(Dnumber),
    CONSTRAINT DEPTSK
        UNIQUE (Dname),
    CONSTRAINT DEPTMGRFK
        FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)
                                ON DELETE SET DEFAULT  ON UPDATE CASCADE);
```

Referential Triggered Action

CREATE TABLE DEPT_LOCATIONS

(. . . ,

PRIMARY KEY (Dnumber, Dlocation),

FOREIGN KEY (Dnumber) **REFERENCES** DEPARTMENT(Dnumber)

ON DELETE CASCADE

ON UPDATE CASCADE);

Trigger

Action

4.2.4 Specifying Constraints on Tuples Using CHECK

```
CHECK (Dept_create_date <= Mgr_start_date);
```

This line is added at the end of the CREATE Table Command

5.4 Schema Change Statements in SQL

Called as **Schema Evolution Commands**

DROP SCHEMA COMPANY

DROP TABLE DEPENDENT

ALTER TABLE COMPANY.EMPLOYEE ADD COLUMN Job VARCHAR(12);

Question Can you use NOT NULL constraint with this new column ?

ALTER TABLE COMPANY.EMPLOYEE DROP COLUMN Address

Queries

Multisets vs Sets

To constrain a table to be a Set, use Key constraint or DISTINCT Command

```
SELECT      <attribute list>
FROM        <table list>
WHERE       <condition>;
```

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

EMPLOYEE

| | | | | | | | | | |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|
| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| | | | |
|-------|----------------|---------|----------------|
| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|-------|----------------|---------|----------------|

DEPT_LOCATIONS

| | |
|----------------|------------------|
| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|

PROJECT

| | | | |
|-------|----------------|-----------|------|
| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-------|----------------|-----------|------|

WORKS_ON

| | | |
|-------------|------------|-------|
| <u>Essn</u> | <u>Pno</u> | Hours |
|-------------|------------|-------|

DEPENDENT

| | | | | |
|-------------|----------------|-----|-------|--------------|
| <u>Essn</u> | Dependent_name | Sex | Bdate | Relationship |
|-------------|----------------|-----|-------|--------------|

Figure 3.5

Schema diagram for the COMPANY relational database schema.

Figure 3.6

One possible database state for the COMPANY relational database schema.

EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|------------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------------|---------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

WORKS_ON

| Essn | Pno | Hours |
|-----------|-----|-------|
| 123456789 | 1 | 32.5 |
| 123456789 | 2 | 7.5 |
| 666884444 | 3 | 40.0 |
| 453453453 | 1 | 20.0 |
| 453453453 | 2 | 20.0 |
| 333445555 | 2 | 10.0 |
| 333445555 | 3 | 10.0 |
| 333445555 | 10 | 10.0 |
| 333445555 | 20 | 10.0 |
| 999887777 | 30 | 30.0 |
| 999887777 | 10 | 10.0 |
| 987987987 | 10 | 35.0 |
| 987987987 | 30 | 5.0 |
| 987654321 | 30 | 20.0 |
| 987654321 | 20 | 15.0 |
| 888665555 | 20 | NULL |

PROJECT

| Pname | Pnumber | Plocation | Dnum |
|-----------------|---------|-----------|------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Stafford | 4 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

DEPENDENT

| Essn | Dependent_name | Sex | Bdate | Relationship |
|-----------|----------------|-----|------------|--------------|
| 333445555 | Alice | F | 1986-04-05 | Daughter |
| 333445555 | Theodore | M | 1983-10-25 | Son |
| 333445555 | Joy | F | 1958-05-03 | Spouse |
| 987654321 | Abner | M | 1942-02-28 | Spouse |
| 123456789 | Michael | M | 1988-01-04 | Son |
| 123456789 | Alice | F | 1988-12-30 | Daughter |
| 123456789 | Elizabeth | F | 1967-05-05 | Spouse |

Query 0. Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

Q0: **SELECT** Bdate, Address
 FROM EMPLOYEE
 WHERE Fname='John' **AND** Minit='B' **AND** Lname='Smith';

Query 1. Retrieve the name and address of all employees who work for the ‘Research’ department.

Q1: **SELECT** Fname, Lname, Address
 FROM EMPLOYEE, DEPARTMENT
 WHERE Dname=‘Research’ **AND** Dnumber=Dno;

Query 2. For every project located in ‘Stafford’, list the project number, the controlling department number, and the department manager’s last name, address, and birth date.

Q2: **SELECT** Pnumber, Dnum, Lname, Address, Bdate
 FROM PROJECT, DEPARTMENT, EMPLOYEE
 WHERE Dnum=Dnumber **AND** Mgr_ssn=Ssn **AND**
 Plocation=‘Stafford’;

(c)

| Pnumber | Dnum | Lname | Address | Bdate |
|---------|------|---------|------------------------|------------|
| 10 | 4 | Wallace | 291Berry, Bellaire, TX | 1941-06-20 |
| 30 | 4 | Wallace | 291Berry, Bellaire, TX | 1941-06-20 |

4.3.2 Ambiguous Attribute Names, Aliasing, Renaming, and Tuple Variables

In SQL, the same name can be used for two (or more) attributes as long as the attributes are in *different relations*. If this is the case, and a multitable query refers to two or more attributes with the same name, we *must qualify* the attribute name with the relation name to prevent ambiguity. This is done by *prefixing* the relation name to the attribute name and separating the two by a period. To illustrate this, suppose that

The ambiguity of attribute names also arises in the case of queries that refer to the same relation twice, as in the following example.

Query 8. For each employee, retrieve the employee's first and last name and the first and last name of his or her immediate supervisor.

Q8: **SELECT** E.Fname, E.Lname, S.Fname, S.Lname
 FROM EMPLOYEE AS E, EMPLOYEE AS S
 WHERE E.Super_ssn=S.Ssn;

One Level
Recursive
Query

E → Subordinates
S → Supervisor

Tuple Variables / Aliases / Copies

You can Rename/Alias without Explicitly using the AS keyword

Q1B: **SELECT** E.Fname, E.LName, E.Address
 FROM EMPLOYEE E, DEPARTMENT D
 WHERE D.DName='Research' **AND** D.Dnumber=E.Dno;

Q9: **SELECT** Ssn
 FROM EMPLOYEE;

Q10: **SELECT** Ssn, Dname
 FROM EMPLOYEE, DEPARTMENT;

Its equivalent to Project Operation in RA

Q1C: **SELECT ***
 FROM EMPLOYEE * Operator, returns all the Attributes
 WHERE Dno=5;

Why SQL Allows Duplicates ?

- Duplicate elimination is an expensive operation. One way to implement it is to sort the tuples first and then eliminate duplicates.
- The user may want to see duplicate tuples in the result of a query.
- When an aggregate function (see Section 5.1.7) is applied to tuples, in most cases we do not want to eliminate duplicates.

Q11: SELECT ~~ALL~~ Salary
 FROM EMPLOYEE;

This ALL is typically not used,
But it is implicit

What if I want to eliminate Duplicates ?

```
Q11A: SELECT      DISTINCT Salary  
        FROM        EMPLOYEE;
```

Set Operations in SQL

No Duplicates
Union
Intersect ...

Duplicates Allowed
Union ALL
Intersect ALL ...

Pattern Matching in SQL

Query 12. Retrieve all employees whose address is in Houston, Texas.

Q12: **SELECT** Fname, Lname
 FROM EMPLOYEE
 WHERE Address **LIKE** '%Houston,TX%';

Query 12A. Find all employees who were born during the 1950s.

Q12: **SELECT** Fname, Lname
 FROM EMPLOYEE
 WHERE Bdate **LIKE** '____5 _____';

What if I want to search for '%' or '_' itself

'AB_CD\%EF' ESCAPE '\'

Usage of Arithmetic operators within Queries

Query 13. Show the resulting salaries if every employee working on the ‘ProductX’ project is given a 10 percent raise.

Q13: **SELECT** E.Fname, E.Lname, 1.1 * E.Salary **AS** Increased_sal
 FROM EMPLOYEE **AS** E, WORKS_ON **AS** W, PROJECT **AS** P
 WHERE E.Ssn=W.Essn **AND** W.Pno=P.Pnumber **AND**
 P.Pname=‘ProductX’;

Between Operator

Query 14. Retrieve all employees in department 5 whose salary is between \$30,000 and \$40,000.

```
Q14:  SELECT      *
        FROM       EMPLOYEE
        WHERE      (Salary BETWEEN 30000 AND 40000) AND Dno = 5;
```

Order By

Query 15. Retrieve a list of employees and the projects they are working on, ordered by department and, within each department, ordered alphabetically by last name, then first name.

```
Q15:  SELECT      D.Dname, E.Lname, E.Fname, P.Pname  
        FROM       DEPARTMENT D, EMPLOYEE E, WORKS_ON W,  
                  PROJECT P  
        WHERE     D.Dnumber= E.Dno AND E.Ssn= W.Essn AND  
                  W.Pno= P.Pnumber  
        ORDER BY  D.Dname, E.Lname, E.Fname;
```

Order By

ORDER BY D.Dname DESC, E.Lname ASC, E.Fname ASC

Structure of a Simple Query

Mandatory Clauses [**SELECT** <attribute list>
 FROM <table list>
 [**WHERE** <condition>]
 [**ORDER BY** <attribute list>];

Optional Clauses

Delete Command

- | | | |
|-------------|--------------------|------------------|
| U4A: | DELETE FROM | EMPLOYEE |
| | WHERE | Lname='Brown'; |
| U4B: | DELETE FROM | EMPLOYEE |
| | WHERE | Ssn='123456789'; |
| U4C: | DELETE FROM | EMPLOYEE |
| | WHERE | Dno=5; |
| U4D: | DELETE FROM | EMPLOYEE; |

Update Command

```
UPDATE      PROJECT  
SET          Plocation = 'Bellaire', Dnum = 5  
WHERE       Pnumber=10;
```

U6:

```
UPDATE      EMPLOYEE  
SET          Salary = Salary * 1.1  
WHERE       Dno = 5;
```

MultiSets in SQL

| First Name | Last Name | Address | City | Age |
|------------|-----------|---------------------|----------|-----|
| Mickey | Mouse | 123 Fantasy Way | Anaheim | 73 |
| Bat | Man | 321 Cavern Ave | Gotham | 54 |
| Wonder | Woman | 987 Truth Way | Paradise | 39 |
| Donald | Duck | 555 Quack Street | Mallard | 65 |
| Bugs | Bunny | 567 Carrot Street | Rascal | 58 |
| Wiley | Coyote | 999 Acme Way | Canyon | 61 |
| Cat | Woman | 234 Purrfect Street | Hairball | 32 |
| Tweety | Bird | 543 | Itotltaw | 28 |

A TABLE in SQL **may** have duplicates, if it has no Primary Key defined.

A TABLE in SQL **can** have duplicates, **If and only if** it has no Primary Key defined.

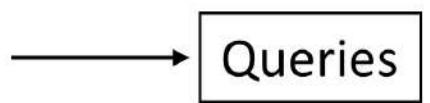
When storing application logs, Caches, or system events, duplicate entries are normal and may even be necessary

```
CREATE TABLE SystemLogs (
    LogID INT AUTO_INCREMENT,
    EventTime TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    EventType VARCHAR(50),
);
```

```
CREATE TABLE API_Cache (
    APIEndpoint VARCHAR(255),
    ResponseData TEXT,
    Timestamp TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

These are highly rare cases, Almost always tables will be created with Primary keys

| First Name | Last Name | Address | City | Age |
|------------|-----------|---------------------|----------|-----|
| Mickey | Mouse | 123 Fantasy Way | Anaheim | 73 |
| Bat | Man | 321 Cavern Ave | Gotham | 54 |
| Wonder | Woman | 987 Truth Way | Paradise | 39 |
| Donald | Duck | 555 Quack Street | Mallard | 65 |
| Bugs | Bunny | 567 Carrot Street | Rascal | 58 |
| Wiley | Coyote | 999 Acme Way | Canyon | 61 |
| Cat | Woman | 234 Purrfect Street | Hairball | 32 |
| Tweety | Bird | 543 | Itotltaw | 28 |



Very High chances for the Presence of
Duplicates.

NULL In SQL

- Unknown
 - Not Available / Unavailable
 - Not Applicable
1. **Unknown value.** A person's date of birth is not known, so it is represented by NULL in the database.
 2. **Unavailable or withheld value.** A person has a home phone but does not want it to be listed, so it is withheld and represented as NULL in the database.
 3. **Not applicable attribute.** An attribute LastCollegeDegree would be NULL for a person who has no college degrees because it does not apply to that person.

NULL In SQL

It is often not possible to determine which of the meanings is intended; for example, a NULL for the home phone of a person can have any of the three meanings. Hence, SQL does not distinguish between the different meanings of NULL.

Three Valued Logic in SQL

Table 5.1 Logical Connectives in Three-Valued Logic

| (a) | AND | TRUE | FALSE | UNKNOWN |
|-----|---------|---------|-------|---------|
| | TRUE | TRUE | FALSE | UNKNOWN |
| | FALSE | FALSE | FALSE | FALSE |
| | UNKNOWN | UNKNOWN | FALSE | UNKNOWN |

| (b) | OR | TRUE | FALSE | UNKNOWN |
|-----|---------|------|---------|---------|
| | TRUE | TRUE | TRUE | TRUE |
| | FALSE | TRUE | FALSE | UNKNOWN |
| | UNKNOWN | TRUE | UNKNOWN | UNKNOWN |

| (c) | NOT | | | |
|-----|---------|---------|--|--|
| | TRUE | FALSE | | |
| | FALSE | TRUE | | |
| | UNKNOWN | UNKNOWN | | |

Three Valued Logic in Practice

```
SELECT * FROM Employee WHERE (Salary > 50000) AND (DeptID = 3);
```

TRUE

Say DeptID = NULL for some rows

AND

UNKNOWN

```
SELECT * FROM Employee WHERE NOT (DeptID = 3);
```

Three Valued Logic in SQL

Query 18. Retrieve the names of all employees who do not have supervisors.

```
SELECT names  
FROM EMPLOYEE  
WHERE Super_ssn = NULL
```

Q18: **SELECT** Fname, Lname
FROM EMPLOYEE
WHERE Super_ssn **IS** NULL;

Nested Queries

Nested Queries

Some queries require that existing values in the database be fetched and then used in a comparison condition. Such queries can be conveniently formulated by using **nested queries**, which are complete select-from-where blocks within the WHERE clause of another query. That other query is called the **outer query**.

Nested Queries

Query 4. Make a list of all project numbers for projects that involve an employee whose last name is ‘Smith’, either as a worker or as a manager of the department that controls the project.

```
Q4A: ( SELECT      DISTINCT Pnumber
        FROM        PROJECT, DEPARTMENT, EMPLOYEE
        WHERE       Dnum=Dnumber AND Mgr_ssn=Ssn
                    AND Lname='Smith' )

        UNION

        ( SELECT      DISTINCT Pnumber
        FROM        PROJECT, WORKS_ON, EMPLOYEE
        WHERE       Pnumber=Pno AND Essn=Ssn
                    AND Lname='Smith' );
```

Nested Queries

Query 4. Make a list of all project numbers for projects that involve an employee whose last name is ‘Smith’, either as a worker or as a manager of the department that controls the project.

Q4A:

```
SELECT      DISTINCT Pnumber
FROM        PROJECT
WHERE       Pnumber IN
            ( SELECT      Pnumber
              FROM        PROJECT, DEPARTMENT, EMPLOYEE
              WHERE       Dnum=Dnumber AND
                          Mgr_ssn=Ssn AND Lname='Smith' )
OR
Pnumber IN
( SELECT      Pno
  FROM        WORKS_ON, EMPLOYEE
  WHERE       Essn=Ssn AND Lname='Smith' );
```

Nested Queries

If a nested query returns a single attribute *and* a single tuple, the query result will be a single (scalar) value. In such cases, it is permissible to use = instead of IN for the comparison operator. In general, the nested query will return a **table** (relation), which is a set or multiset of tuples.

- **Nested** subquery runs only once for the entire nesting (outer) query. It does not contain any reference to the outer query row.
- **Correlated** subquery runs once for each row selected by the outer query. It contains a reference to a value from the row selected by the outer query.

| # | Correlated Subquery | Non-Correlated Subquery |
|-------------------|---|--|
| Execution | Triggered for each row in the outer query | Executed only once, regardless of outer query rows |
| Dependency | Relies on values from the outer query | Independent of the outer query |
| Efficiency | Less efficient for large datasets | More efficient for large datasets |

Nested Queries

```
SELECT      DISTINCT Essn  
FROM        WORKS_ON  
WHERE       (Pno, Hours) IN ( SELECT      Pno, Hours  
                           FROM        WORKS_ON  
                           WHERE      Essn='123456789' );  
  
  
Notice the parenthesis
```

Comparing Single Value to a Set of Values

In addition to the `IN` operator, a number of other comparison operators can be used to compare a single value v (typically an attribute name) to a set or multiset V (typically a nested query). The `= ANY` (or `= SOME`) operator returns `TRUE` if the value v is equal to *some value* in the set V and is hence equivalent to `IN`. The two keywords `ANY` and `SOME` have the same effect. Other operators that can be combined with `ANY` (or `SOME`) include `>`, `>=`, `<`, `<=`, and `<>`. The keyword `ALL` can also be combined with each of these operators. For example, the comparison condition ($v > \text{ALL } V$) returns `TRUE` if the value v is greater than *all* the values in the set (or multiset) V .

Comparison operators in Nested Queries

```
SELECT      Lname, Fname  
FROM        EMPLOYEE  
WHERE       Salary > ALL    ( SELECT      Salary  
                           FROM        EMPLOYEE  
                           WHERE       Dno=5 );
```

Query 16. Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```
SELECT      E.Fname, E.Lname
FROM        EMPLOYEE AS E
WHERE       E.Ssn IN  ( SELECT      Essn
                      FROM        DEPENDENT AS D
                      WHERE       E.Fname=D.Dependent_name
                      AND E.Sex=D.Sex );
```

```
SELECT      E.Fname, E.Lname
FROM        EMPLOYEE AS E
WHERE       E.Ssn IN  ( SELECT      Essn
                        FROM        DEPENDENT AS D
                        WHERE       E.Fname=D.Dependent_name
                        AND E.Sex=D.Sex );
```

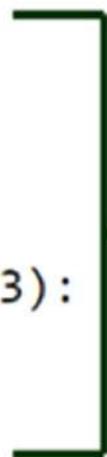
In the nested query of Q16, we must qualify E.Sex because it refers to the Sex attribute of EMPLOYEE from the outer query, and DEPENDENT also has an attribute called Sex. If there were any unqualified references to Sex in the nested query, they would refer to the Sex attribute of DEPENDENT. However, we would not *have to* qualify the attributes Fname and Ssn of EMPLOYEE if they appeared in the nested query because the DEPENDENT relation does not have attributes called Fname and Ssn, so there is no ambiguity.

5.1.3 Correlated Nested Queries

Whenever a condition in the WHERE clause of a nested query references some attribute of a relation declared in the outer query, the two queries are said to be **correlated**.

The Inner Query is evaluated once for each row of the outer query

```
for i in range(2):
    print(i)
    for j in range(10,13):
        print(j)
```



Outer loop

Inner loop

Output

↓

0
10
11
12
1
10
11
12

Printed by inner loop

Printed by outer loop

A diagram showing the output sequence. The numbers 0, 10, 11, 12, 1, 10, 11, and 12 are listed vertically. Brackets on the left group 0, 10, 11, and 12 as "Printed by inner loop". Brackets on the right group 1, 10, 11, and 12 as "Printed by outer loop". An arrow points down from the word "Output".

```
SELECT      E.Fname, E.Lname
FROM        EMPLOYEE AS E
WHERE       E.Ssn IN  ( SELECT      Essn
                        FROM        DEPENDENT AS D
                        WHERE       E.Fname=D.Dependent_name
                        AND E.Sex=D.Sex );
```

For each EMPLOYEE tuple:

retrieve the Essn values for all DEPENDENT tuples with the same sex and name as that EMPLOYEE tuple;

If the Ssn value of the EMPLOYEE tuple is in the result of the nested query, then select that EMPLOYEE tuple.

In general, a query written with nested select-from-where blocks and using the = or IN comparison operators can *always* be expressed as a single block query.

```
SELECT          E.Fname, E.Lname  
FROM           EMPLOYEE AS E, DEPENDENT AS D  
WHERE          E.Ssn=D.Essn AND E.Sex=D.Sex  
                  AND E.Fname=D.Dependent_name;
```

5.1.4 The EXISTS and UNIQUE Functions in SQL

The **EXISTS** function in SQL is used to check whether the result of a correlated nested query is empty (contains no tuples) or not.

```
SELECT      E.Fname, E.Lname  
FROM        EMPLOYEE AS E  
WHERE       EXISTS ( SELECT *  
                      FROM    DEPENDENT AS D  
                      WHERE   E.Ssn=D.Essn AND E.Sex=D.Sex  
                             AND E.Fname=D.Dependent_name);
```

For each Employee:
If there exist atleast one dependent with
Same sex and same name, then
Return that employee

If at least one tuple EXISTS in the result of the nested query, then select that EMPLOYEE tuple

```
SELECT      Fname, Lname
FROM        EMPLOYEE
WHERE       NOT EXISTS ( SELECT      *
                           FROM        DEPENDENT
                           WHERE       Ssn=Essn );
```

Query 6. Retrieve the names of employees who have no dependents.

If and only if “NO” Tuple EXISTS in the result of the nested query, then select that EMPLOYEE tuple

Query 7. List the names of managers who have at least one dependent.

```
SELECT      Fname, Lname  
FROM        EMPLOYEE  
WHERE       EXISTS ( SELECT *  
                      FROM   DEPENDENT  
                      WHERE  Ssn=Essn )  
  
AND  
EXISTS ( SELECT *  
          FROM  DEPARTMENT  
          WHERE  Ssn=Mgr_ssn );
```

Selects Employees
With Dependents

Selects Employees
Who are managers

For Each Employee tuple:
You return that tuple
If they have both
dependents
And
Also they are managers

Query 7. List the names of managers who have at least one dependent.

```
SELECT DISTINCT E.Fname, E.Lname  
FROM EMPLOYEE E, DEPARTMENT D, DEPENDENT DEP  
WHERE E.Ssn = D.Mgr_ssn  
AND E.Ssn = DEP.Essn;
```

Query 3. Find the names of employees who work on *all* the projects controlled by department number 5.

```
SELECT      Fname, Lname
FROM        EMPLOYEE
WHERE       NOT EXISTS ( ( SELECT      Pnumber
                           FROM        PROJECT
                           WHERE       Dnum=5)
EXCEPT      ( SELECT      Pno
                           FROM        WORKS_ON
                           WHERE       Ssn=Essn) );
```

Not Correlated with the Outer Query

Selects all projects that the particular employee being considered works on

For each employee:
IF (First Subquery – Second Subquery) = Empty
It means that the employee works on all the projects and is therefore selected

```
SELECT      E.Fname, E.Lname  
FROM        EMPLOYEE AS E  
WHERE       E.Ssn IN  ( SELECT      Essn  
                      FROM        DEPENDENT AS D  
                      WHERE       E.Fname=D.Dependent_name  
                      AND E.Sex=D.Sex );
```

For each EMPLOYEE tuple:

 Retrieve the Essn values for all DEPENDENT tuples with the same sex and name as that EMPLOYEE tuple;

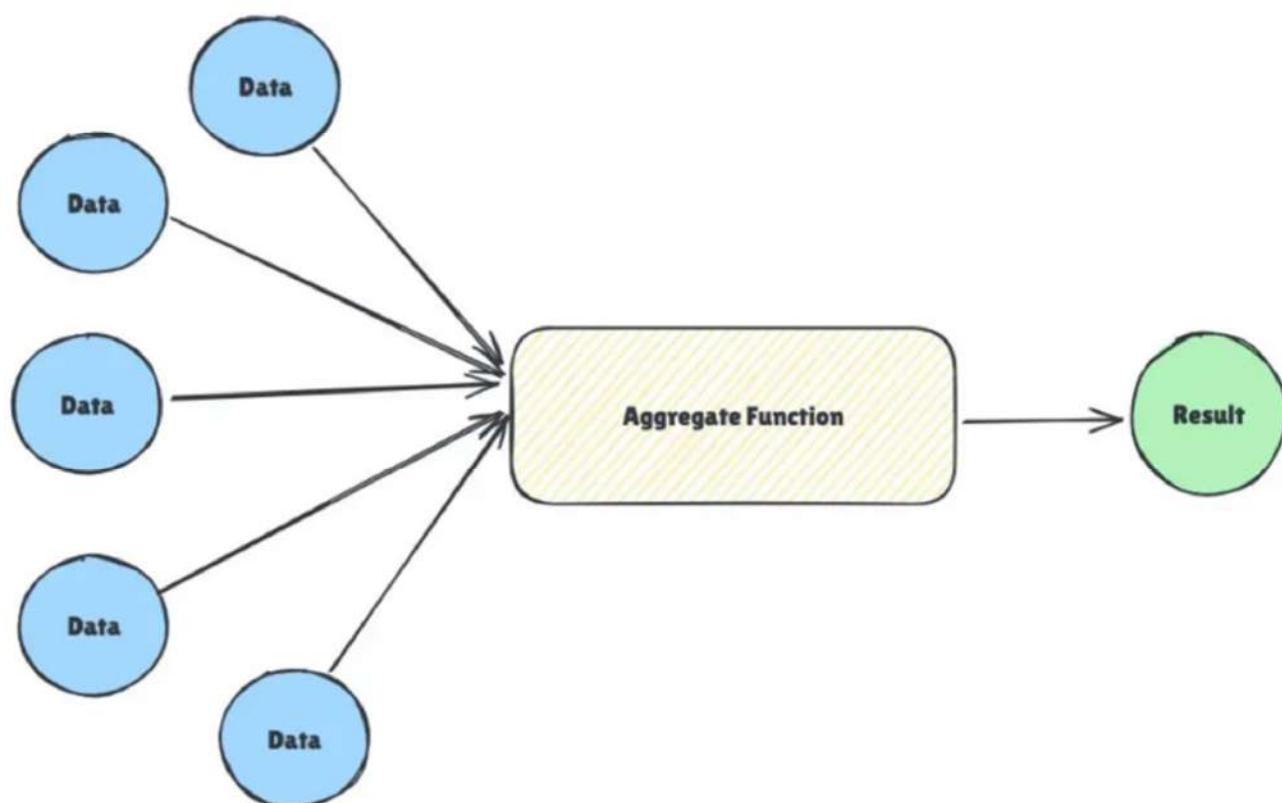
 If the Ssn value of the EMPLOYEE tuple matches with the any of the Essn Value, then return the F.name and L.Name

```
SELECT      E.Fname, E.Lname  
FROM        EMPLOYEE AS E, DEPENDENT AS D  
WHERE       E.Ssn=D.Essn AND E.Sex=D.Sex  
              AND E.Fname=D.Dependent_name;
```

In general, a query written with nested select-from-where blocks and using the = or IN comparison operators can *always* be expressed as a single block query. For exam-

5.1.7 Aggregate Functions in SQL

Aggregate functions are used to summarize information from multiple tuples into a single-tuple summary.



| Function Name | Meaning | Example |
|------------------|---|-------------|
| SUM(column name) | Total sum of the values in a numeric column | SUM(salary) |
| AVG(column name) | Average of the values in a column | AVG(salary) |
| MAX(column name) | Largest value in a column | MAX(salary) |
| MIN(column name) | Smallest value in a column | MIN(salary) |
| COUNT(*) | Count of the number of rows selected | COUNT(*) |

Query 20. Find the sum of the salaries of all employees of the ‘Research’ department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
SELECT      SUM (Salary), MAX (Salary), MIN (Salary), AVG (Salary)  
FROM        (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)  
WHERE       Dname=‘Research’;
```

Count the number of distinct salary values in the database.

```
SELECT          COUNT (DISTINCT Salary)  
FROM            EMPLOYEE;
```

What if someone's salary is NULL, will it be counted ?

In **general**, NULL values are discarded when aggregate functions are applied to a particular column (attribute).

```
SELECT      COUNT (*)  
FROM        EMPLOYEE;
```

```
SELECT      COUNT (*)  
FROM        EMPLOYEE, DEPARTMENT  
WHERE       DNO=DNUMBER AND DNAME='Research';
```

```
SELECT      Lname, Fname  
FROM        EMPLOYEE  
WHERE       (SELECT COUNT (*)  
                  FROM DEPENDENT  
                  WHERE Ssn=Essn ) >= 2;
```

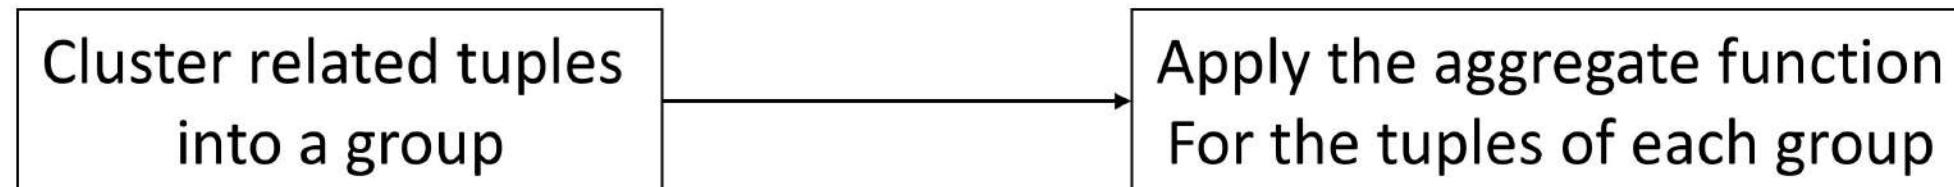
This correlated nested query counts the number of dependents that each employee has;
if this is greater than or equal to two, the employee tuple is selected

Consider queries like:

Find average salary of each department.

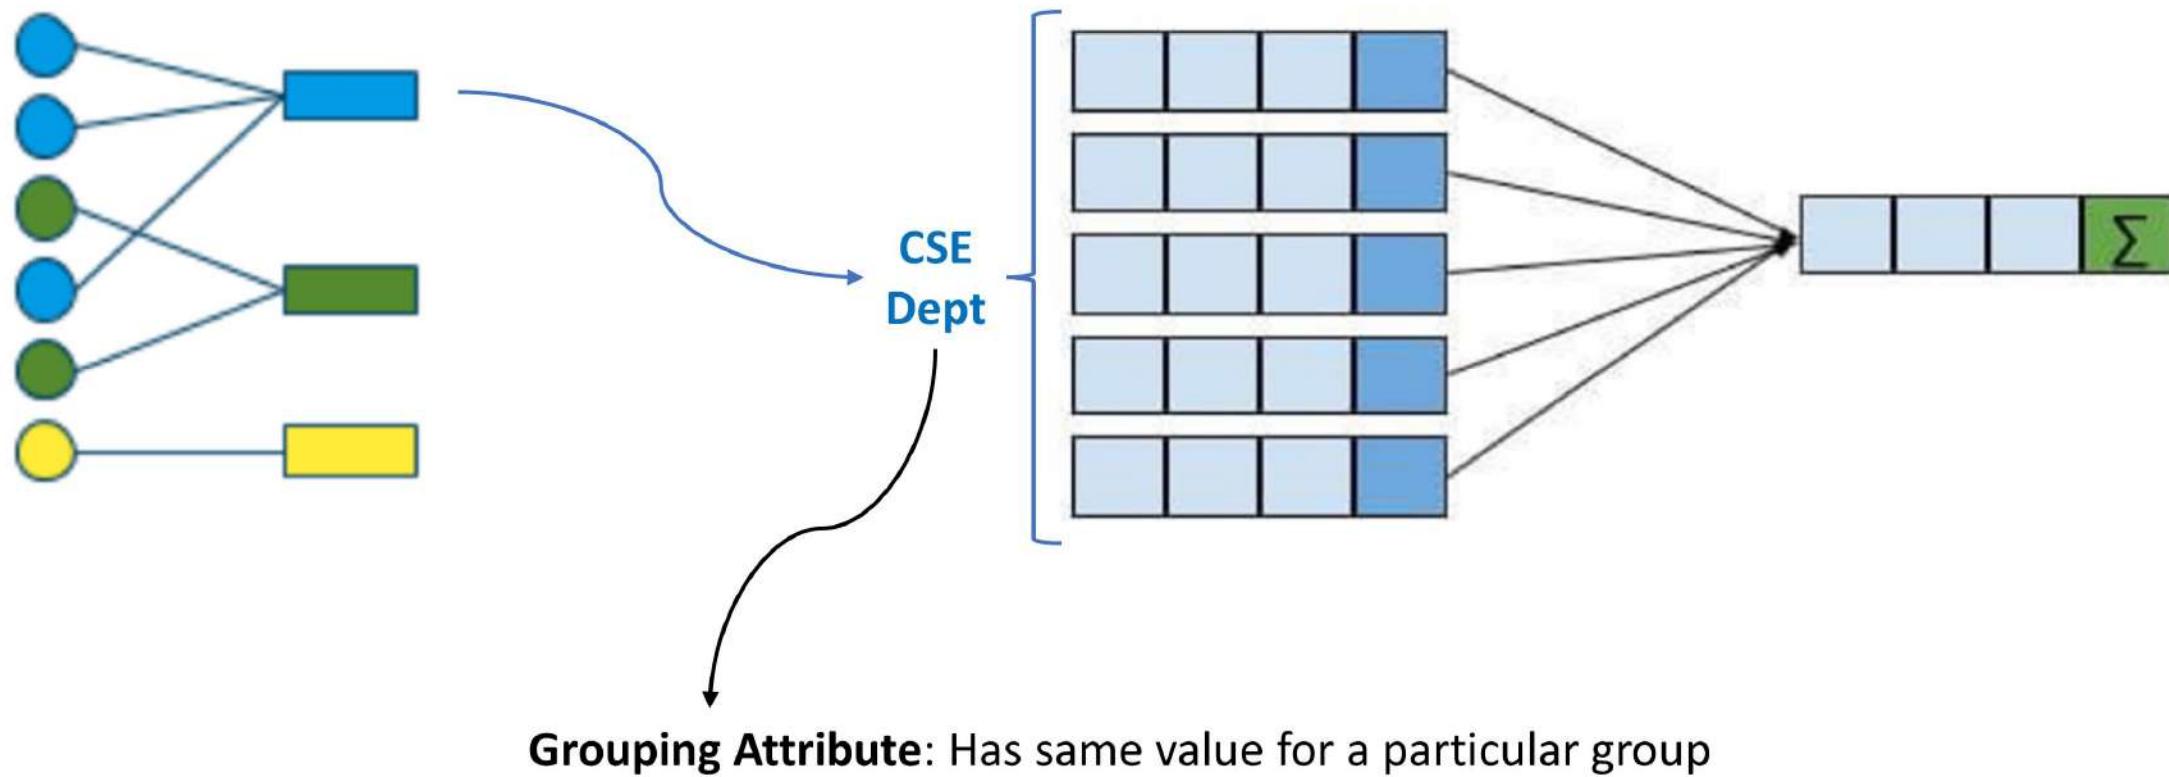
Find the number of employees in each project.

Find the number of employees managed by each manager.



Cluster related tuples
into a group

Apply the aggregate function
For the tuples of each group



Query 24. For each department, retrieve the department number, the number of employees in the department, and their average salary.

Q24: **SELECT** Dno, **COUNT (*)**, **AVG** (Salary)
 FROM EMPLOYEE
 GROUP BY Dno;

Note that: Select clause includes only the Grouping Attribute and the Aggregate functions to be applied on each Group of tuples

Q24: **SELECT** Dno, COUNT (*), AVG (Salary)
FROM EMPLOYEE
GROUP BY Dno;

(a)

| Fname | Minit | Lname | Ssn | ... | Salary | Super_ssn | Dno | | Dno | Count (*) | Avg (Salary) |
|----------|-------|---------|-----------|-----|--------|-----------|-----|--|-----|-----------|--------------|
| John | B | Smith | 123456789 | ... | 30000 | 333445555 | 5 | | 5 | 4 | 33250 |
| Franklin | T | Wong | 333445555 | | 40000 | 888665555 | 5 | | 4 | 3 | 31000 |
| Ramesh | K | Narayan | 666884444 | | 38000 | 333445555 | 5 | | 1 | 1 | 55000 |
| Joyce | A | English | 453453453 | | 25000 | 333445555 | 5 | | | | |
| Alicia | J | Zelaya | 999887777 | | 25000 | 987654321 | 4 | | | | |
| Jennifer | S | Wallace | 987654321 | | 43000 | 888665555 | 4 | | | | |
| Ahmad | V | Jabbar | 987987987 | | 25000 | 987654321 | 4 | | | | |
| James | E | Bong | 888665555 | | 55000 | NULL | 1 | | | | |

Result of Q24

Grouping EMPLOYEE tuples by the value of Dno

What if there were some employees whose Dno = NULL

Query 25. For each project, retrieve the project number, the project name, and the number of employees who work on that project.

Q25: **SELECT** Pnumber, Pname, **COUNT (*)**
 FROM PROJECT, WORKS_ON
 WHERE Pnumber=Pno
 GROUP BY Pnumber, Pname;

Query 25. For each project, retrieve the project number, the project name, and the number of employees who work on that project.



Say I'm adding one more condition to this query
I want only projects with more than two employees to appear in the result



Query 26. For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

Q26: **SELECT** Pnumber, Pname, **COUNT** (*)
 FROM PROJECT, WORKS_ON
 WHERE Pnumber=Pno
 GROUP BY Pnumber, Pname
 HAVING **COUNT** (*) > 2;

WHERE helps to choose tuples to which the aggregate functions are applied.

HAVING clause serves to choose whole groups.

| WHERE Clause | HAVING Clause |
|---|---|
| Applicable without GROUP BY clause | Cannot be used without GROUP BY Clause |
| Can be used with SELECT, UPDATE, DELETE statement. | Can only be used with SELECT statement |
| Cannot contain aggregate function | Can contain aggregate function |
| Used to filter the records from the table based on the specified condition. | Used to filter record from the groups based on the specified condition. |

```
SELECT <attribute and function list>
FROM <table list>
[ WHERE <condition> ]
[ GROUP BY <grouping attribute(s)> ]
[ HAVING <group condition> ]
[ ORDER BY <attribute list> ];
```

IN Vs Exists

IN is used when comparing a column against a list of values or a subquery result set.

It checks if a value is present in the specified list or subquery result.

The subquery returns a list of values that the main query checks against

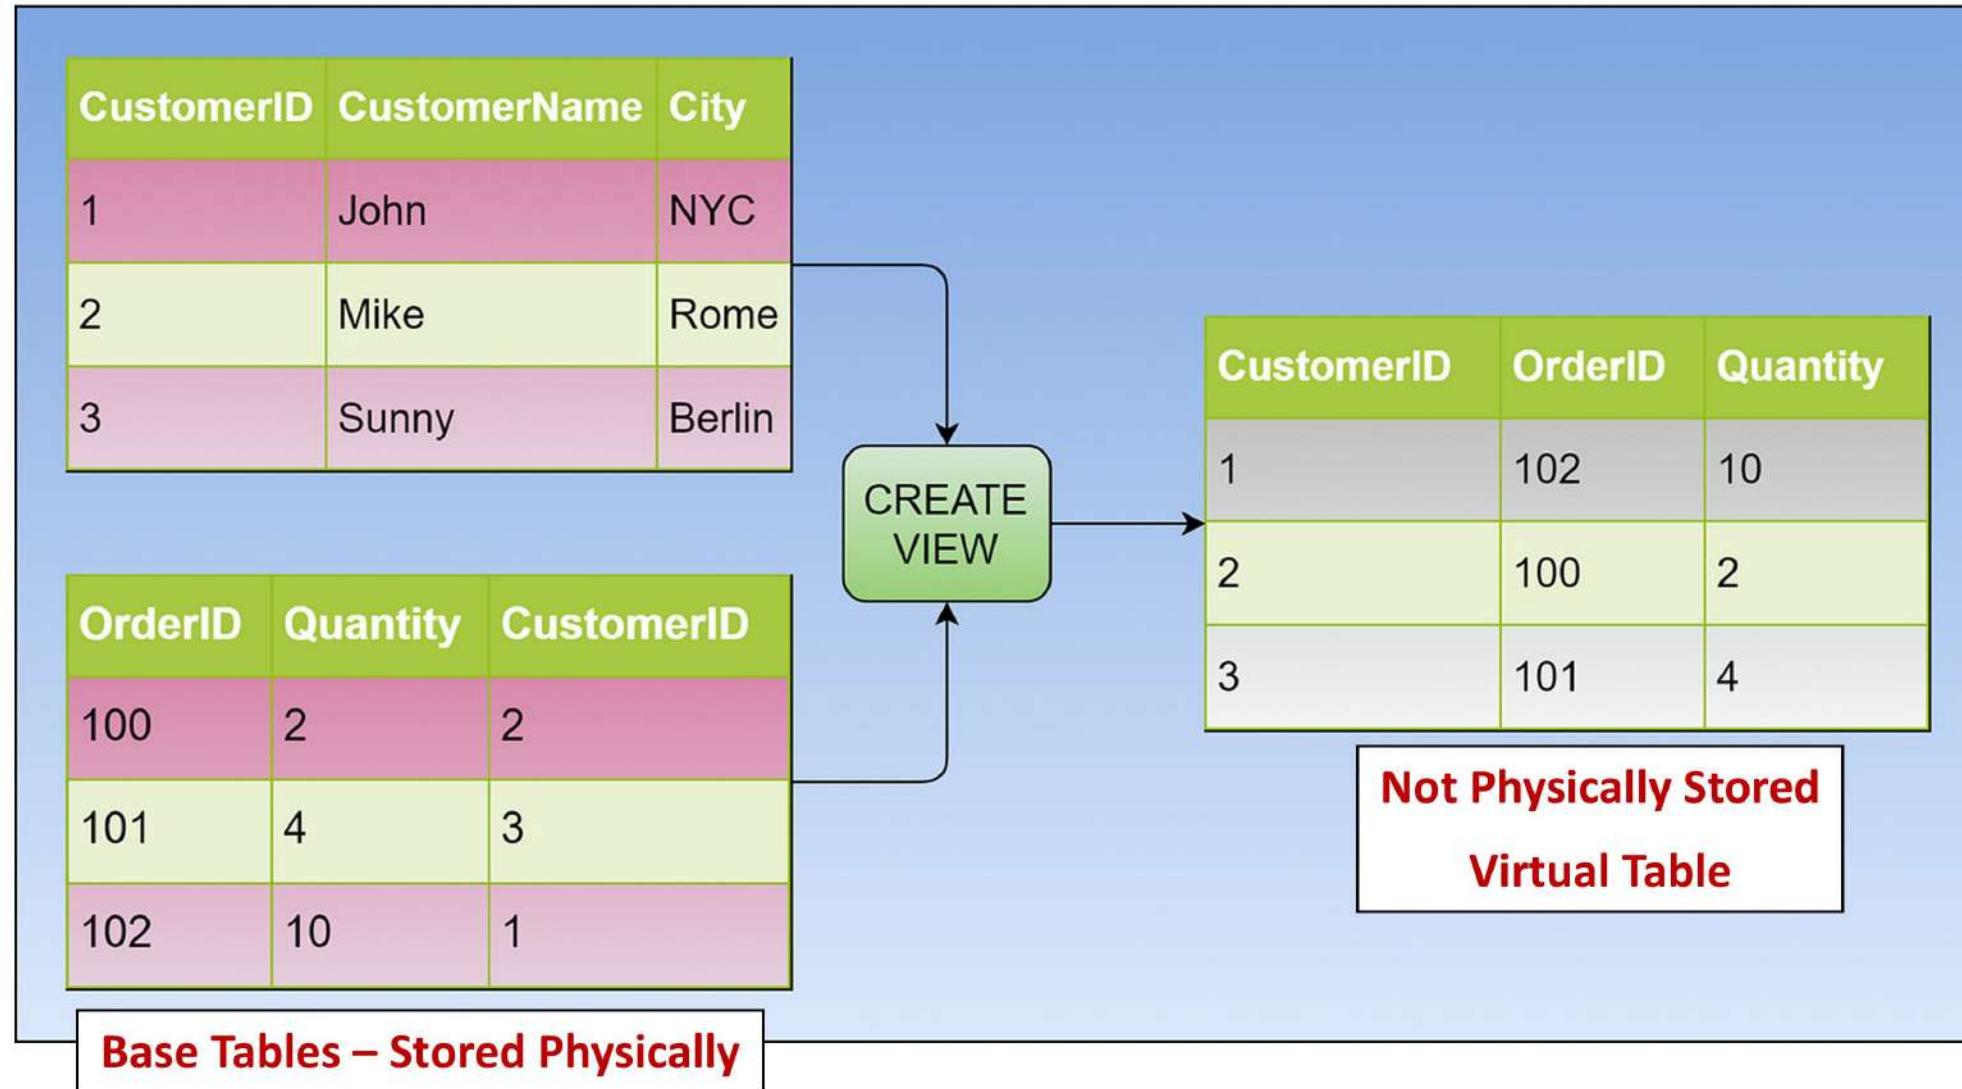
EXISTS checks whether the subquery returns any rows.

It does not return values, only TRUE or FALSE.

Stops execution as soon as a match is found (more efficient for large datasets).

Best suited when checking if a condition is met rather than retrieving values.

| Feature | IN | EXISTS |
|--------------|-----------------------------------|-----------------------------|
| Works on | List of values | Subquery returning rows |
| Returns | Values from subquery | TRUE or FALSE |
| Performance | Slower for large data | Faster for large data |
| Stops early? | No, checks all values | Yes, stops at first match |
| Use case | When checking for specific values | When checking for existence |



V1: **CREATE VIEW** WORKS_ON1
 AS SELECT Fname, Lname, Pname, Hours
 FROM EMPLOYEE, PROJECT, WORKS_ON
 WHERE Ssn=Essn **AND** Pno=Pnumber;

VIEWS are not generally subjected to Updation.

VIEWS can be queried just like any other table.

SELECT Fname, Lname
FROM WORKS_ON1
WHERE Pname='ProductX';

```
V2:   CREATE VIEW DEPT_INFO(Dept_name, No_of_emps, Total_sal)
      AS SELECT Dname, COUNT (*), SUM (Salary)
              FROM DEPARTMENT, EMPLOYEE
             WHERE Dnumber=Dno
        GROUP BY Dname;
```



DEPT_INFO

| Dept_name | No_of_emps | Total_sal |
|-----------|------------|-----------|
|-----------|------------|-----------|

Querying the VIEWS

QV1: **SELECT** Fname, Lname
FROM WORKS_ON1
WHERE Pname='ProductX';  **SELECT** Fname, Lname
EMPLOYEE, PROJECT, WORKS_ON
Ssn=Essn **AND** Pno=Pnumber
AND Pname='ProductX';

VIEW Materialization

Updating the VIEWS

- A view with a single defining table is updatable if the view attributes contain the primary key of the base relation, as well as all attributes with the NOT NULL constraint *that do not have* default values specified.
- Views defined on multiple tables using joins are generally not updatable.
- Views defined using grouping and aggregate functions are not updatable.

V1A: DROP VIEW WORKS_ON1;

Assertions – Specifying General Constraints

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK ( NOT EXISTS ( SELECT      *
                      FROM        EMPLOYEE E, EMPLOYEE M,
                                  DEPARTMENT D
                     WHERE      E.Salary>M.Salary
                               AND E.Dno=D.Dnumber
                               AND D.Mgr_ssN=M.Ssn ) );
```

The salary of an employee must not be greater than the salary of the manager of the department that the employee works for.

The basic technique for writing such assertions is to specify a query that selects any tuples that violate the desired condition.

By including this query inside a NOT EXISTS clause, the assertion will specify that the result of this query must be empty so that the condition will always be TRUE

Thus, the assertion is violated if the result of the query is not empty

SQL triggers are a critical feature in **database management systems (DBMS)** that provide automatic execution of a set of SQL statements when specific database events, such as **INSERT**, **UPDATE**, or **DELETE** operations, occur.

Triggers are commonly used to maintain **data integrity**, **track changes**, and **enforce business rules** automatically, without needing manual input.

For example, a trigger can be invoked when a row is inserted into a specified table or when specific table columns are updated

```
create trigger [trigger_name]  
[before | after]  
{insert | update | delete}  
on [table_name]  
FOR EACH ROW  
BEGIN  
END;
```

Key Terms

- **trigger_name:** The name of the trigger to be created.
- **BEFORE | AFTER:** Specifies whether the trigger is fired **before** or **after** the triggering event (INSERT, UPDATE, DELETE).
- **{INSERT | UPDATE | DELETE}:** Specifies the operation that will activate the trigger.
- **table_name:** The name of the table the trigger is associated with.
- **FOR EACH ROW:** Indicates that the trigger is row-level, meaning it executes once for each affected row.
- **trigger_body:** The SQL statements to be executed when the trigger is fired.