

IMPLEMENTING NETWORK SOFTWARE

10 September
2025

Dr Noor Mahammad Sk

Implementing Network Software

2

- Network architectures and protocol specifications are essential things to define a good blue print
- The number of computers connected to the Internet has roughly doubled every 12 to 18 months since 1981
- Number of Internet users are increasing exponentially over the year
- Great success of Internet contributed by
 - ▣ A good architecture
 - ▣ Major contribution is its functionality is provided by the software running in general purpose computers
 - The significance of this is that new functionality can be added readily with “just a small matter of programming”

Application Programming Interface (Sockets)

3

- Since most network protocols are implemented in software (especially those high in the protocol stack)
- All computers implement their network protocols as part of the operating system
- The interface that OS provides to its networking subsystem
 - ▣ Is called as the network Application Programming Interface (API)
- Each operating system is free to define its own network API
 - ▣ Over time certain of these APIs have become widely supported

Application Programming Interface (Sockets)

4

- ❑ The *socket interface* originally provided by the Berkeley distribution of Unix
- ❑ It is supported in virtually all popular operating systems
- ❑ The advantage of industry-wide support for a single API is that applications can be easily ported from one OS to another
 - ▣ Developers can easily write applications for multiple OSs
- ❑ The network application programs typically interact with many parts of the OS other than the network

Application Programming Interface (Sockets)

5

- Two systems support the same network API does not mean that their file system, process, or graphic interfaces are the same
- Each protocol provides a certain set of services
- The API provides a *syntax* by which those services can be invoked in this particular OS
- The socket is the point where a local application process attaches to the network
- The interface defines operations for
 - ▣ creating a socket,
 - ▣ attaching the socket to network
 - ▣ Sending/receiving message through the socket, and closing the socket

Application Programming Interface (Sockets)

6

- First step to create a socket
 - ▣ *int socket(int domain, int type, int protocol)*
- The domain argument specifies the protocol family that is going to be used
 - ▣ PF_INET denotes the internet family
 - ▣ PF_UNIX denotes the Unix pipe facility
 - ▣ PF_PACKET denotes direct access to network interface(i.e., it bypasses the TCP/IP protocol stack)
- The type argument indicates the semantic of the communication
 - ▣ SOCK_STREAM denotes a byte stream
 - ▣ SOCK_DGRAM is an alternative that denotes a message-oriented service, such as that provided by UDP

Application Programming Interface (Sockets)

7

- ❑ On a server machine the application process performs a passive open
 - ▣ The server says that it is prepared to accept connections,
 - ▣ but it does not actually establish a connection
 - ▣ The server does the connection by invoking the following three operations
 - `int bind(int socket, struct sockaddr *address, int addr_len)`
 - `int listen(int socket, int backlog)`
 - `int accept(int socket, struct sockaddr *address, int *addr_len)`
- ❑ The bind operation: binds the newly created socket to the specified address

Application Programming Interface (Sockets)

8

- ❑ The listen operation defines how many connections can be pending on the specified *socket*
- ❑ Accept operation carries out the passive open
 - ▣ It is a blocking operation that does not return until a remote participant has established a connection
 - ▣ When it does complete, it returns a new socket that corresponds to this just established connection
- ❑ The address argument contains the *remote* participant's address

Application Programming Interface (Sockets)

9

- ❑ On the client machine, the application process performs an *active* open;
 - ▣ It says who it wants to communicate
 - ▣ *int connect(int socket, struct sockaddr *address, int addr_len)*
 - ▣ This process does not return until TCP has successfully established a connection at which time the application is free to begin sending data
- ❑ In practice, the client usually, specifies only the remote participant's address and lets the system fill in the local information


Application Programming Interface (Sockets)

10

- Server usually listens for messages on a well-known port, a client typically does not care which port it uses for itself
 - ▣ The OS simply selects an unused one
- Once connection is established, the application processes invoke the following two operations to send and receive data
 - ▣ *int send(int socket, char *message, int msg_len, int flags)*
 - Sends the given message over the specified socket
 - ▣ *int recv(int socket, char *buffer, int buf_len, int flags)*
 - Receives a message from the specified socket into the given buffer
 - ▣ Both the operations take a set of flags that control certain details of the operation

End System: Computer on the 'Net

11



Also known as a “host”...

Clients and Servers


12

□ Client program

- ▣ Running on end host
- ▣ Requests service
- ▣ E.g., Web browser

□ Server program

- ▣ Running on end host
- ▣ Provides service
- ▣ E.g., Web server



Clients Are Not Necessarily Human


13

- Example: Web crawler (or spider)
 - ▣ Automated client program
 - ▣ Tries to discover & download many Web pages
 - ▣ Forms the basis of search engines like Google
- Spider client
 - ▣ Start with a base list of popular Web sites
 - ▣ Download the Web pages
 - ▣ Parse the HTML files to extract hypertext links
 - ▣ Download these Web pages, too
 - ▣ And repeat, and repeat, and repeat...

Client-Server Communication

14

- Client “sometimes on”
 - Initiates a request to the server when interested
 - E.g., Web browser on your laptop or cell phone
 - Doesn’t communicate directly with other clients
 - Needs to know server’s address
- Server is “always on”
 - Services requests from many client hosts
 - E.g., Web server for the www.iiitdm.ac.in Web site
 - Doesn’t initiate contact with the clients
 - Needs fixed, known address



Peer-to-Peer Communication

15

- No always-on server at the center of it all
 - ▣ Hosts can come and go, and change addresses
 - ▣ Hosts may have a different address each time
- Example: peer-to-peer file sharing
 - ▣ Any host can request files, send files, query to find a file's location, respond to queries, ...
 - ▣ Scalability by harnessing millions of peers
 - ▣ Each peer acting as both a client and server

Client and Server Processes

16

- Program vs. process
 - ▣ Program: collection of code
 - ▣ Process: a running program on a host
- Communication between processes
 - ▣ Same end host: inter-process communication
 - Governed by the operating system on the end host
 - ▣ Different end hosts: exchanging messages
 - Governed by the network protocols
- Client and server processes
 - ▣ Client process: process that initiates communication
 - ▣ Server process: process that waits to be contacted

Delivering the Data: Division of Labor


17

- Network
 - ▣ Deliver data packet to the destination host
 - ▣ Based on the destination IP address
- Operating system
 - ▣ Deliver data to the destination socket
 - ▣ Based on the destination port number (e.g., 80)
- Application
 - ▣ Read data from and write data to the socket
 - ▣ Interpret the data (e.g., render a Web page)

Socket: End Point of Communication

18

- Sending message from one process to another
 - ▣ Message must traverse the underlying network
- Process sends and receives through a “socket”
 - ▣ In essence, the doorway leading in/out of the house
- Socket as an Application Programming Interface
 - ▣ Supports the creation of network applications



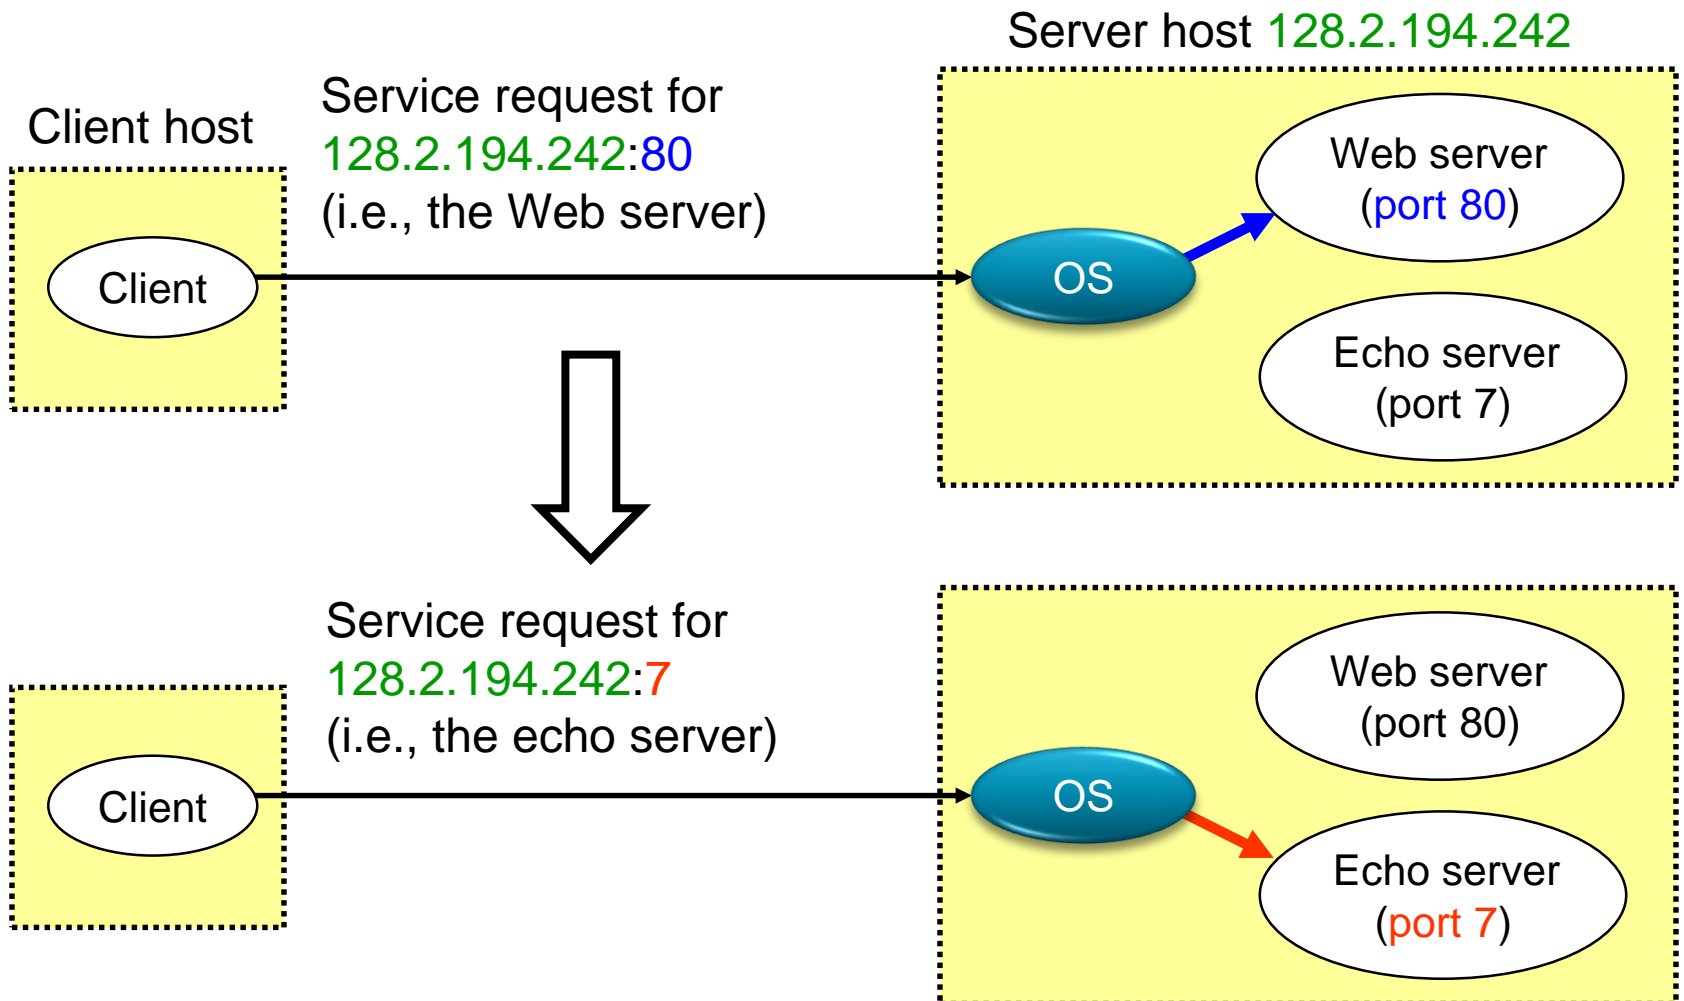
Identifying the Receiving Process

19

- Sending process must identify the receiver
 - ▣ The receiving end host machine
 - ▣ The specific socket in a process on that machine
- Receiving host
 - ▣ Destination address that uniquely identifies the host
 - ▣ An IP address is a 32-bit quantity
- Receiving socket
 - ▣ Host may be running many different processes
 - ▣ Destination port that uniquely identifies the socket
 - ▣ A port number is a 16-bit quantity

Using Ports to Identify Services

20



Knowing What Port Number To Use

21

- Popular applications have well-known ports
 - ▣ E.g., port 80 for Web and port 25 for e-mail
 - ▣ See <http://www.iana.org/assignments/port-numbers>
- Well-known vs. ephemeral ports
 - ▣ Server has a well-known port (e.g., port 80)
 - Between 0 and 1023 (requires root to use)
 - ▣ Client picks an unused ephemeral (i.e., temporary) port
 - Between 1024 and 65535
- Uniquely identifying traffic between the hosts
 - ▣ Two IP addresses and two port numbers
 - ▣ Underlying transport protocol (e.g., TCP or UDP)

Port Numbers are Unique per Host

22

- Port number uniquely identifies the socket
 - ▣ Cannot use same port number twice with same address
 - ▣ Otherwise, the OS can't demultiplex packets correctly
- Operating system enforces uniqueness
 - ▣ OS keeps track of which port numbers are in use
 - ▣ Doesn't let the second program use the port number
- Example: two Web servers running on a machine
 - ▣ They cannot both use port "80", the standard port #
 - ▣ So, the second one might use a non-standard port #
 - ▣ E.g., <http://www.iiitdm.ac.in:8080>

UNIX Socket API

23

- Socket interface
 - ▣ Originally provided in Berkeley UNIX
 - ▣ Later adopted by all popular operating systems
 - ▣ Simplifies porting applications to different OSES
- In UNIX, everything is like a file
 - ▣ All input is like reading a file
 - ▣ All output is like writing a file
 - ▣ File is represented by an integer file descriptor
- API implemented as system calls
 - ▣ E.g., connect, read, write, close, ...

Typical Client Program

24

- Prepare to communicate
 - ▣ Create a socket
 - ▣ Determine server address and port number
 - ▣ Initiate the connection to the server
- Exchange data with the server
 - ▣ Write data to the socket
 - ▣ Read data from the socket
 - ▣ Do stuff with the data (e.g., render a Web page)
- Close the socket

Servers Differ From Clients

25

- ❑ Passive open
 - ▣ Prepare to accept connections
 - ▣ ... but don't actually establish
 - ▣ ... until hearing from a client
- ❑ Hearing from multiple clients
 - ▣ Allowing a backlog of waiting clients
 - ▣ ... in case several try to communicate at once
- ❑ Create a socket for each client
 - ▣ Upon accepting a new client
 - ▣ ... create a *new* socket for the communication



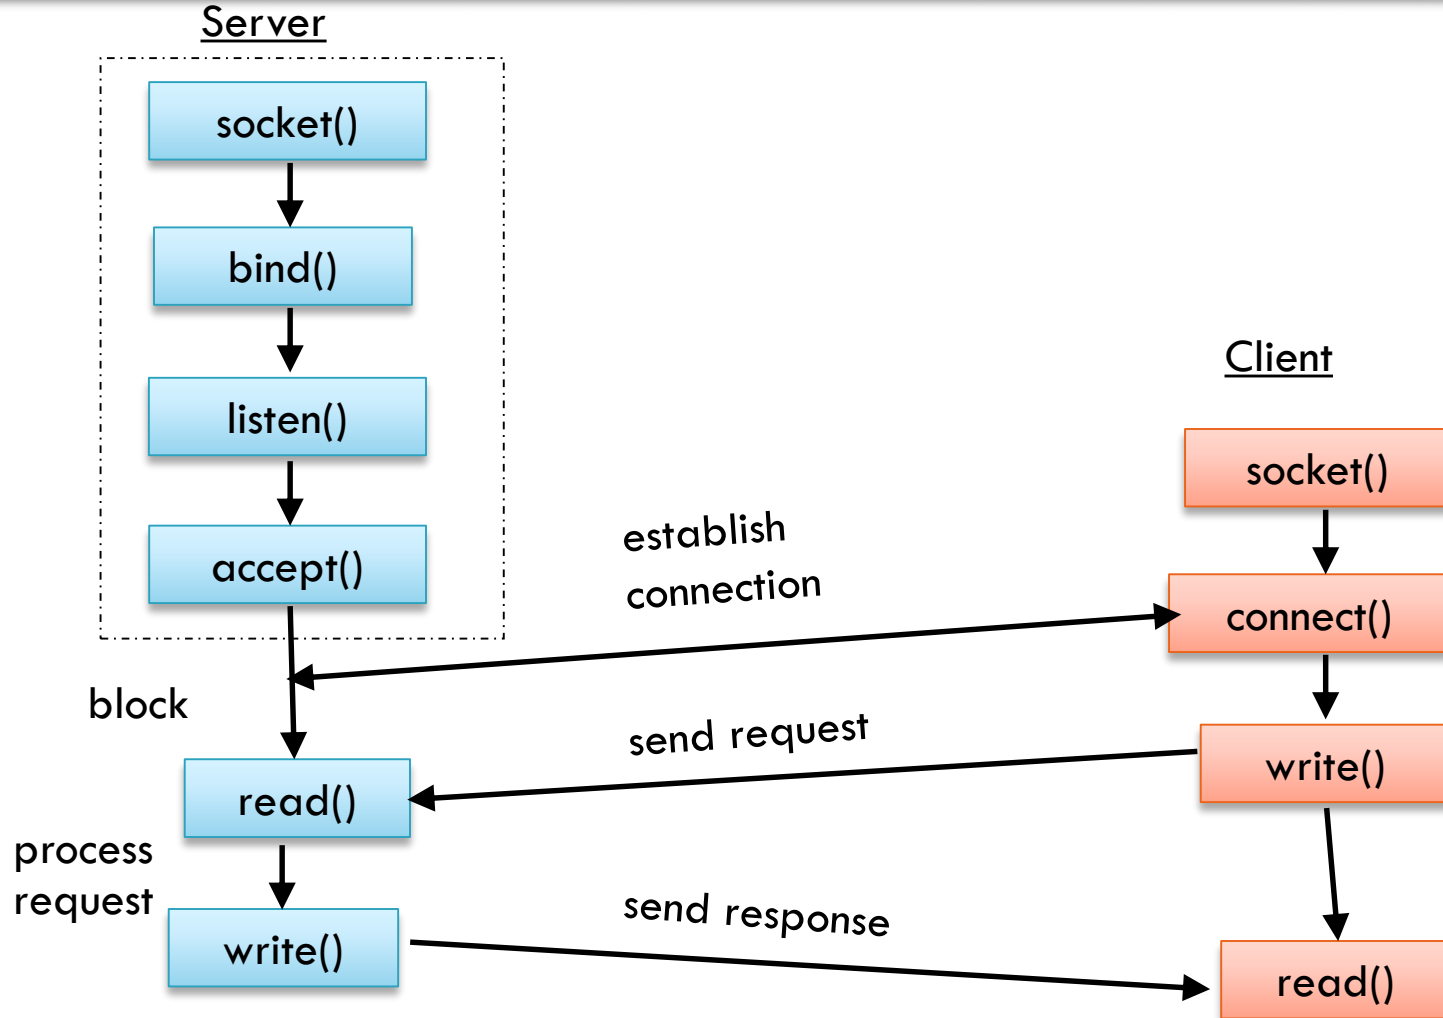
Typical Server Program

26

- ❑ Prepare to communicate
 - ▣ Create a socket
 - ▣ Associate local address and port with the socket
- ❑ Wait to hear from a client (passive open)
 - ▣ Indicate how many clients-in-waiting to permit
 - ▣ Accept an incoming connection from a client
- ❑ Exchange data with the client over new socket
 - ▣ Receive data from the socket
 - ▣ Do stuff to handle the request (e.g., get a file)
 - ▣ Send data to the socket
 - ▣ Close the socket
- ❑ Repeat with the next connection request

Putting it All Together

27



Client Creating a Socket: `socket()`

28

- ❑ Creating a socket
 - ▣ *`int socket(int domain, int type, int protocol)`*
 - ▣ Returns a file descriptor (or handle) for the socket
 - ▣ Originally designed to support any protocol suite
- ❑ Domain: protocol family
 - ▣ `PF_INET` for the Internet (IPv4)
- ❑ Type: semantics of the communication
 - ▣ `SOCK_STREAM`: reliable byte stream (TCP)
 - ▣ `SOCK_DGRAM`: message-oriented service (UDP)
- ❑ Protocol: specific protocol
 - ▣ `UNSPEC`: unspecified

Client: Learning Server Address/Port

29

- Server typically known by name and service
 - ▣ E.g., “www.cnn.com” and “http”
- Need to translate into IP address and port #
 - ▣ E.g., “64.236.16.20” and “80”
- Translating the server’s name to an address
 - ▣ ***struct hostent *gethostbyname(char *name)***
 - ▣ Argument: host name (e.g., “www.cnn.com”)
 - ▣ Returns a structure that includes the host address
- Identifying the service’s port number
 - ▣ ***struct servent***
 - ▣ ****getservbyname(char *name, char *proto)***
 - ▣ Arguments: service (e.g., “ftp”) and protocol (e.g., “tcp”)
 - ▣ Static config in `/etc/services`

Client: Connecting Socket to the Server

30

- ❑ Client contacts the server to establish connection
 - ▣ Associate the socket with the server address/port
 - ▣ Acquire a local port number (assigned by the OS)
 - ▣ Request connection to server, who hopefully accepts

- ❑ Establishing the connection
 - ▣ *int connect (int sockfd, struct sockaddr *server_address, socketlen_t addrlen)*
 - ▣ Arguments: socket descriptor, server address, and address size
 - ▣ Returns 0 on success, and -1 if an error occurs

Client: Sending Data

31

- Sending data
 - ▣ ***ssize_t write (int sockfd, void *buf, size_t len)***
 - ▣ Arguments: socket descriptor, pointer to buffer of data to send, and length of the buffer
 - ▣ Returns the number of bytes written, and -1 on error

Client: Receiving Data

32

- Receiving data
 - ▣ *ssize_t read(int sockfd, void *buf, size_t len)*
 - ▣ Arguments: socket descriptor, pointer to buffer to place the data, size of the buffer
 - ▣ Returns the number of characters read (where 0 implies “end of file”), and -1 on error

- Closing the socket
 - ▣ *int close(int sockfd)*

Server: Server Preparing its Socket

33

- Server creates a socket and binds address/port
 - ▣ Server creates a socket, just like the client does
 - ▣ Server associates the socket with the port number (and hopefully no other process is already using it!)
 - ▣ Choose port “0” and let kernel assign ephemeral port
- Create a socket
 - ▣ ***int socket (int domain, int type, int protocol)***
- Bind socket to the local address and port number
 - ▣ ***int bind (int sockfd, struct sockaddr *my_addr, socklen_t addrlen)***
 - ▣ Arguments: sockfd, server address, address length
 - ▣ Returns 0 on success, and -1 if an error occurs

Server: Allowing Clients to Wait

34

- ❑ Many client requests may arrive
 - ▣ Server cannot handle them all at the same time
 - ▣ Server could reject the requests, or let them wait
- ❑ Define how many connections can be pending
 - ▣ ***int listen(int sockfd, int backlog)***
 - ▣ Arguments: socket descriptor and acceptable backlog
 - ▣ Returns a 0 on success, and -1 on error
- ❑ What if too many clients arrive?
 - ▣ Some requests don't get through
 - ▣ The Internet makes no promises...
 - ▣ And the client can always try again

Server: Accepting Client Connection

35

- Now all the server can do is wait...
 - ▣ Waits for connection request to arrive
 - ▣ Blocking until the request arrives
 - ▣ And then accepting the new request

- Accept a new connection from a client
 - ▣ ***int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen)***
 - ▣ Arguments: sockfd, structure that will provide client address and port, and length of the structure
 - ▣ Returns descriptor of socket for this new connection

Server: One Request at a Time?

36

- ❑ Serializing requests is inefficient
 - ▣ Server can process just one request at a time
 - ▣ All other clients must wait until previous one is done
- ❑ May need to time share the server machine
 - ▣ Alternate between servicing different requests
 - Do a little work on one request, then switch when you are waiting for some other resource (e.g., reading file from disk)
 - “Nonblocking I/O”
 - ▣ Or, use a different process/thread for each request
 - Allow OS to share the CPU(s) across processes
 - ▣ Or, some hybrid of these two approaches

Client and Server: Cleaning House

37

- Once the connection is open
 - ▣ Both sides can read and write
 - ▣ Two unidirectional streams of data
 - ▣ In practice, client writes first, and server reads
 - ▣ ... then server writes, and client reads, and so on
- Closing down the connection
 - ▣ Either side can close the connection
 - ▣ ... using the `close()` system call
- What about the data still “in flight”
 - ▣ Data in flight still reaches the other end
 - ▣ So, server can `close()` before client finishes reading


One Annoying Thing: Byte Order

38

- ❑ Hosts differ in how they store data
 - ▣ E.g., four-byte number (byte3, byte2, byte1, byte0)
- ❑ Little endian (“little end comes first”): Intel x86’s
 - ▣ Low-order byte stored at the lowest memory location
 - ▣ Byte0, byte1, byte2, byte3
- ❑ Big endian (“big end comes first”)
 - ▣ High-order byte stored at lowest memory location
 - ▣ Byte3, byte2, byte1, byte 0
- ❑ Makes it more difficult to write portable code
 - ▣ Client may be big or little endian machine
 - ▣ Server may be big or little endian machine

Endian Example: Where is the Byte?

39



8 bits memory

16 bits Memory

32 bits Memory

Little-
Endian

1000	78
1001	
1002	
1003	

	+1	+0
1000		78
1002		
1004		
1006		

	+3	+2	+1	+0
1000				78
1004				
1008				
100C				

Big-
Endian

	+0	+1
1000	78	
1001		
1002		
1003		

	+0	+1
1000	78	
1002		
1004		
1006		

	+0	+1	+2	+3
1000	78			
1004				
1008				
100C				

IP is Big Endian

40

- ❑ But, what byte order is used “on the wire”
 - ▣ That is, what do the network protocol use?
- ❑ The Internet Protocols picked one convention
 - ▣ IP is big endian (aka “network byte order”)
- ❑ Writing portable code require conversion
 - ▣ Use htons() and htonl() to convert to network byte order
 - ▣ Use ntohs() and ntohl() to convert to host order
- ❑ Hides details of what kind of machine you’re on
 - ▣ Use the system calls when sending/receiving data structures longer than one byte

Using htonl and htons

41

```
int sockfd = // connected SOCK_STREAM
u_int32_t my_val = 1234;
u_int16_t my_xtra = 16;

u_short bufsize = sizeof (struct data_t);
char *buf = New char[bufsize];
bzero (buf,  bufsize);

struct data_t *dat = (struct data_t *) buf;
dat->value = htonl (my_val);
dat->xtra  = htons (my_xtra);

int rc = write (sockfd, buf, bufsize);
```

42

Implementing Network Software

Protocol Implementation Issues

Protocol Implementation Issues

43

- A high level protocol interacts with a low-level protocol
 - ▣ Example: TCP needs an interface to send outgoing messages to IP
 - IP needs to be able to deliver incoming messages to TCP
- Process Model
 - ▣ Most operating systems provide an abstraction called a *process*, or alternatively, a *thread*
 - ▣ Each process runs largely independently of other processes
 - ▣ The OS is responsible for making sure that resources
 - Such as address space and CPU cycles, are allocated to all the current processes

Process Model

44

- The process abstraction makes it fairly straightforward to have a lot of things executing concurrently on one machine
 - ▣ Example: each user application might execute in its own process, and various things inside the OS might execute as other processes
- When the OS stops one process from executing on the CPU and starts up another one, we call the change a *context switch*


Network Process Models

45


- Two models
- Process-per-protocol model
 - ▣ Each protocol is implemented by a separate process
 - ▣ As a message moves up or down the protocol stack
 - It is passed from one process/protocol to another
 - Example: process that implements protocol i processes the message, then passes it to protocol $i-1$ and so on
 - ▣ Host OS provides interprocess communication
 - To support one process/protocol passes a message to the next process/protocol

Network Process Models

46



Process-per-Protocol



Process-per-Message

Network Process Models

47

- Process-per-message
 - ▣ Treats each protocol as a static piece of code and associates the processes with the message
 - ▣ When message arrives from the network, the OS dispatches a process that it makes responsible for the message as it moves up the protocol graph
 - ▣ At each level, the procedure that implements that protocol is invoked
 - Which eventually results in the procedure for the next protocol being invoked and so on
 - ▣ In both directions, the protocol graph is traversed in a sequence of procedure calls

Process-per-Model vs process-per-message

48

- Process-per-Model
 - ▣ A context switch is required at each level of the protocol graph – typically time consuming
 - ▣ Requires a context switch at each level
- process-per-message
 - ▣ The process-per-message model is generally more efficient
 - A procedure call is an order of magnitude more efficient than a context switch on most computers
 - ▣ A procedure call per level


Message Buffers

49

- The application process provides the buffer that contains the outbound message
 - ▣ When calling ***send***
- Similarly it provides the buffer into which an incoming message is copied
 - ▣ when invoking the ***receive*** operation
- This forces the topmost protocol to copy the message from the application's buffer into a network buffer, and vice versa

Message Buffers

50



Copying incoming/outgoing messages between application buffer and network buffer

Message Buffers –Drawback


51

- Copying data from one buffer to another is one of the most expensive things a protocol implementation can do
 - ▣ Memory speed is slower when compared to processors
- Most of networks defines an abstract data type for messages that is shared by all protocols in the protocol graph
- This abstraction permits
 - ▣ Message to be passed up and down the protocol graph without copying
 - ▣ It also provides a copy free ways of manipulation of messages
 - Such as adding and stripping headers
 - Fragmenting large messages into a set of small messages
 - Reassembling a collection of small messages into a single large message

Message Buffers

52

- The exact form of message abstraction differs from OS to OS
- In generally involves a linked list of pointers to message buffers



Example message data Structure

PERFORMANCE

10 September
2025

Dr Noor Mahammad Sk

Performance

54

- The effectiveness of computations distributed over the network often depends directly on the efficiency with which the network delivers the computation's data
- Network performance is measured in two fundamental ways
 - ▣ Bandwidth (throughput)
 - ▣ Latency (delay)

55

Performance

Bandwidth and Latency

Bandwidth and Latency


56

- Bandwidth: the number of bits that can be transmitted over the network in a certain period of time
 - ▣ Example: 10 millions bits per seconds
 - It is able to deliver 10 million bits every second
 - ▣ On 10 Mbps network, it takes 0.1 microsecond (μs) to transmit each bit
- Latency: corresponding to how long it takes a message to travel from one end of a network to the other
 - ▣ Latency is measured strictly in terms of time


Bandwidth

57

- Bits transmitted at a particular bandwidth can be regarded as having some width



Bits transmitted at 1mbps (each bit 1 μ s wide)



Bits transmitted at 2mbps (each bit 0.5 μ s wide)

Latency

58

- **Round-trip time (RTT)**: the time it takes to send a message from one end of a network to the other and back
- Latency depends of three parameters
- Propagation Delay
- **Transmit Delay**: the amount of time it take to transmit a unit of data
- **Queue Delays**: since packets switches generally need to store packets for some time before forwarding them on an outbound link
- ***Latency = Propagation + Transmit + Queue***
- ***Propagation = Distance/SpeedOfLight***
- ***Transmit = Size/Bandwidth***

Latency

59

- **Distance**: Length of the wire over which the data will travel
- **SpeeOfLight**: effective speed of light over that wire
- **Size**: size of the packet
- **Bandwidth**: bandwidth at which the packet is transmitted

Performance

60

- Bandwidth and latency combine to define the performance characteristics of a given link or channel
- Latency dominates bandwidth for some applications
 - ▣ A client that sends a 1-byte message to a server and receives a 1-byte message in return is latency bound
 - ▣ The application performs much differently on different channels
 - A transcontinental channel with a 100ms RTT
 - An across the room channel with a 1ms RTT
 - ▣ The channel is 1Mbps or 100Mbps is relatively insignificant
 - To transmit a byte on a above will take 8us and 0.08us time.

Performance

61

- ❑ A digital library program that is being asked to fetch a 25MB image
- ❑ The more bandwidth that is available, the faster it will be able to return the image to the user
- ❑ The bandwidth of the channel dominates performance

Delay x Bandwidth Product

62

- A channel a pair of processes as a hollow pipe
 - The latency corresponds to the length of the pipe
 - Bandwidth gives the diameter of the pipe
 - The delay x bandwidth product gives the volume of the pipe
 - For example a transcontinental channel with a one-way latency of 50ms and a bandwidth of 45 Mbps is able to hold
 - $50 \times 10^{-3} \text{ sec} \times 45 \times 10^6 \text{ bits/sec}$
 - $2.25 \times 10^6 \text{ bits}$



Delay x Bandwidth Product

63

Link Type	Bandwidth (typical)	Distance (typical)	Round-trip Delay	Delay x BW
Dail – up	56 Kbps	10Km	87 μ s	5bits
Wireless LAN	54 Mbps	50m	0.33 μ s	18 bits
Satellite	45 Mbps	35,000 Km	230 ms	10Mb
Cross-country fiber	10Gbps	4,000 km	40 ms	400 Mb

- The delay x bandwidth product is important to know then constructing high-performance networks
- Because it corresponds to how many bits the sender must transmit before the first bit arrives at the receiver

Delay x Bandwidth Product

64

- Relative importance of bandwidth and latency depends on application
 - ▣ For large file transfer, bandwidth is critical
 - ▣ For small messages (HTTP, NFS, etc.), latency is critical
 - ▣ Variance in latency (jitter) can also affect some applications (*e.g.*, audio/video conferencing)

Delay x Bandwidth Product

65

- How many bits the sender must transmit before the first bit arrives at the receiver if the sender keeps the pipe full
- Takes another one-way latency to receive a response from the receiver
- If the sender does not fill the pipe—send a whole delay \times bandwidth product's worth of data before it stops to wait for a signal—the sender will not fully utilize the network


Delay x Bandwidth Product

66

- Infinite bandwidth
 - ▣ RTT dominates
 - ▣ $\text{Throughput} = \text{TransferSize} / \text{TransferTime}$
 - ▣ $\text{TransferTime} = \text{RTT} + (1/\text{Bandwidth}) \times \text{TransferSize}$
- Its all relative
 - ▣ 1-MB file to 1-Gbps link looks like a 1-KB packet to 1-Mbps link

Relationship between Bandwidth and Latency

67



A 1-MB file would fill the 1-Mbps link 80 times,
but only fill the 1-Gbps link 1/12 of one time

Throughput of Network

68

- $\text{Throughput} = \text{TransferSize} / \text{TransferTime}$
- $\text{TransferTime} = \text{RTT} + (1/\text{Bandwidth}) \times \text{TransferSize}$

Summary

69

- We have identified what we expect from a computer network
- We have defined a layered architecture for computer network that will serve as a blueprint for our design
- We have discussed the socket interface which will be used by applications for invoking the services of the network subsystem
- We have discussed two performance metrics using which we can analyze the performance of computer networks

THANK YOU!!

10 September
2025

Noor Mahammad Sk